

## 1 Answers

This pdf document contains solutions for questions asked as part of the “Genome Annotation” module.

### 1.1 Answers for ‘Comparing Reference Genomes’:

**Question: Based on your work during the previous assembly module, can you think of a reason why assembly might not be perfect?**

There are many possible answers here (refer to slides 3 & 4 from the assembly presentation), including:

1. Repetitive sequences.
2. Low sequencing depth.
3. Errors in reads

**Question: Is there an obvious issue with our assembly?**

Yes, the 3’ end of our assembly is missing!

**Question: Why do you think both ends of the reference genome align to the same part of our assembled genome?**

Both ends of a chromosome have telomeres, which contain identical sequence.

**Questions: What do you think the green segments represent in this image?**

Inversions – pieces of the genome that are in the opposite direction when compared between isolates.

**Question: Why is the red line not centered in the plot and moves up or down?**

Some pieces of the genome are missing from one isolate to another, meaning they align at different places.

### 1.2 Answers for ‘Identifying Repetitive DNA’:

Below are the repetitive sequences in BOLD RED:

```
TATAAAATACAATATAATATAAACGACGAACAGATATGAAAGTGTTAGAACTAGACATACCA
TTTTTCTGTGAAAAATACTTCAAGCTGTAGTATTATTATTATTGCGCTGCTTAGATGTAGT
```

**Question: Why do the sections “Retroelements” and “DNA transposons” all have zeros?**

*Plasmodium* species do not have transposons.

**Question: Approximately what proportion of our genome assembly was masked?**

We can look at the file `PB.contigs.polished.reheader.fasta.tbl` using a command like `less`, `more`, or `cat`. You should see a table similar to the following image:

	number of elements*	length occupied	percentage of sequence
Retroelements	0	0 bp	0.00 %
SINEs:	0	0 bp	0.00 %
Penelope	0	0 bp	0.00 %
LINEs:	0	0 bp	0.00 %
CRE/SLACS	0	0 bp	0.00 %
L2/CR1/Rex	0	0 bp	0.00 %
R1/LOA/Jockey	0	0 bp	0.00 %
R2/R4/NeSL	0	0 bp	0.00 %
RTE/Bov-B	0	0 bp	0.00 %
L1/CIN4	0	0 bp	0.00 %
LTR elements:	0	0 bp	0.00 %
BEL/Pao	0	0 bp	0.00 %
Ty1/Copia	0	0 bp	0.00 %
Gypsy/DIRS1	0	0 bp	0.00 %
Retroviral	0	0 bp	0.00 %
DNA transposons	0	0 bp	0.00 %
hobo-Activator	0	0 bp	0.00 %
Tc1-IS630-Pogo	0	0 bp	0.00 %
En-Spm	0	0 bp	0.00 %
MuDR-IS905	0	0 bp	0.00 %
PiggyBac	0	0 bp	0.00 %
Tourist/Harbinger	0	0 bp	0.00 %
Other (Mirage, P-element, Transib)	0	0 bp	0.00 %
Rolling-circles	0	0 bp	0.00 %
Unclassified:	0	0 bp	0.00 %
Total interspersed repeats:		0 bp	0.00 %
Small RNA:	0	0 bp	0.00 %
Satellites:	0	0 bp	0.00 %
Simple repeats:	3629	178057 bp	14.62 %
Low complexity:	527	30591 bp	2.51 %

From here you can see that ~17% of the genome was masked (Simple repeats + Low complexity).

### 1.3 Answers for 'Finding Genes':

**Question: How many exons does the gene “1\_g” have?**

Looking at the output of the head command you can see that ‘1\_g’ has 4 exons. To get this answer, you can also try the more expert-level command:

```
grep "\"1_g\"" genemark.gtf | grep exon | wc -l
```

This command should print 4 on the command line. The two \" in the command above tell grep to look for quotation marks as part of the search. The \ is what is known as an “escape character”.

**Question: Can you think of a simple LINUX command to figure out how many genes GeneMark-ES identified?**

There are multiple solutions to this question – ask your instructor if how you found the solution is acceptable. One possible solution is to use the tool awk like below:

```
[ ]: awk '{print $10}' genemark.gtf | uniq | wc -l
```

This command works by:

1. Printing the 10th column (awk '{print \$10}'; “\$10” is the column with gene\_id).
2. *PIPING* the output of awk to the next command, uniq
3. uniq finds all *unique* values
4. *PIPING* the output of uniq to the next command, wc
5. wc -l tells you how many lines are printed by the previous command.

To see each step in this command, try running each of the above without the next pipe, like so:

1. awk '{print \$10}' genemark.gtf
2. awk '{print \$10}' genemark.gtf | uniq
3. awk '{print \$10}' genemark.gtf | uniq | wc -l

**Question: How many genes did GeneMark find?**

You can use the above command to find out. The answer should be close to 300 Genes.

**Question: What is the part of the command 2> /dev/null actually doing?**

Directing (also known as “piping”) anything to /dev/null means you are directing any output to an empty file that is not actually record

**Question: How many genes are in our final set of possible genes (bonafide.gb)?**

This requires you to come up with another command by using some abstract problem-solving. If you look at the bonafide.gb file like so:

```
more bonafide.gb
```

You may notice that for each gene record, there is a single “LOCUS” line. We can use that to find out how many genes were found:

```
grep -c LOCUS bonafide.gb
```

the -c flag counts the number of times LOCUS was found within the file bonafide.gb

**Question: What do you think a limitation of using just 1 chromosome to train our gene finder is?**

The main limitation is the number of genes that we can use for training. As we are only using 1 chromosome to train AUGUSTUS, there are only ~300 genes that can be used for training. If we use all chromosomes, that means we can use all possible genes in the entire genome (~5,300 genes).

**Question: Can you figure out how many genes each approach found?**

We can use a similar command as the previous question where we found the number of genes in `bonafide.gb`:

```
grep -c 'start gene' PB.contigs.default.gff
```

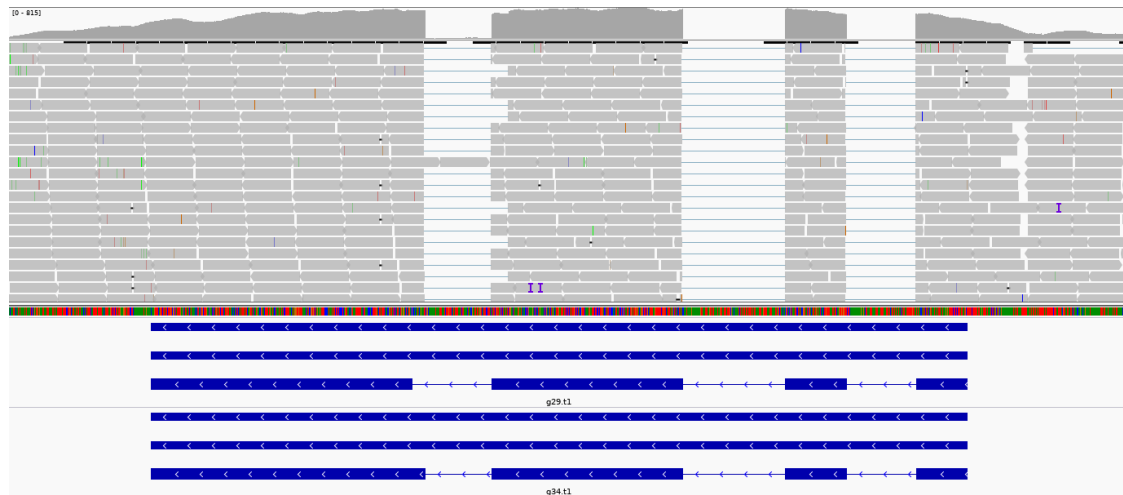
Since each gff file contains one 'start gene' entry per found gene, we can use `grep` to count the total number of identified genes. This should give a value of approximately 300-320 genes for each file.

**Question: We identified protein coding genes, but can you think of any other types of annotations we could find with Augustus?**

If you read the AUGUSTUS paper, you can also identify: UnTranslated Regions (UTRs) and Processed pseudogenes.

**Question: How many exons does this gene have?**

See the below figure:



This gene has 4 exons.

**Question: How do the predictions for your model versus the default model compare?**

The default model uses an alternative splice site for one of the exons.

**Question: What are both predictions missing, and why do you think that is?**

They are both missing UTRs. That is because we did not tell AUGUSTUS to search for them.

**Question: How many exons does this gene have?**

It has 8 exons.

**Question: How many introns?**

It has 7 introns.

**Question: How can you extract the same information for another gene by modifying the above command? Report the command and result here.**

You can just change “g54” to any other gene number and then modify the number given to -A to change the number of lines. An example for another gene (g53) is:

```
grep -A 15 "start gene g53" PB.contigs.gff
```

## 1.4 Answers for 'Using Comparative Genomes to Identify Genes':

**Question: What is difficult about this alignment?**

It has an intron.

**Question: Did you notice something at the end of the alignment that was not in the protein sequence?**

The stop codon “TAG”.

**Question: What was difficult in this example?**

There is a mutation in the DNA sequence that caused a different amino acid to be in the sequence. The sequence changed from GCA to GGA.

**Question: Do you think this is an issue, or is there something biology-related going on?**

Changes in DNA sequence occur all the time between species. The change in amino acid could reflect some selective advantage for a given species or be neutral and confer no evolutionary advantage.

**Question: What do you think the \* character represents?**

It is the symbol the stop codon.

**Question: What is the gene that we identified in IGV?**

Myosin B

**Question: Can you name a function of this gene and how did you get the answer?**

Myosins are motor proteins that are used for motility via binding of actin chains in eukaryotes. There are a number of resources where this information can be found, including the UCSC genome browser, Uniprot, and *even* Wikipedia – ;D.